Bhavya Priya N

Data Engineer — AWS Data Platforms

Bengaluru, India | +91 9445202321 | bhavyapriya127@gmail.com
LinkedIn: linkedin.com/in/bhavyapriya | GitHub: github.com/BhavyapriyaN | Portfolio: bhavyapriyan.github.io/portfolio

Professional Summary

Data Engineer with 3 years of experience building cloud ELT pipelines and analytics platforms with Python, SQL and AWS (Redshift, Glue, Athena, S3, Lambda, DynamoDB). Experience across billing data migration, incremental ingestion, data modeling, governance and automation on billing data from 600+ AWS accounts, including 50+ Glue jobs, a 75% pipeline runtime reduction, and approximately $39,000 per month in savings from an automated WorkSpaces cleanup workflow.

Professional Experience

Datagrokr Analytics Pvt. Ltd., Bengaluru, India — Oct. 2023 – Present
Data Engineer

- Built a serverless Amazon WorkSpaces cleanup workflow with 3 Lambda functions after the move to Windows remote desktop: flagged WorkSpaces inactive for 3+ months, alerted owners, raised ServiceNow termination tickets after 14 days without action, and tracked ticket status daily; saved approximately $39,000 per month.
- Migrated the production AWS Cost and Usage Report (CUR) platform from 1.0 to 2.0 (billing data for 600+ AWS accounts), redesigning Redshift schemas and reloading 15 months of history (800M rows); handled nested attributes with Redshift SUPER and cut query runtime from 18–20 hours to 7 hours through distribution and sort key tuning.
- Orchestrated 50+ AWS Glue jobs via dependency-based Glue Workflows in Step Functions-initiated ELT pipelines loading CUR data into Redshift in about 40 minutes daily; monitored with CloudWatch, applied VACUUM maintenance, validated net amortized costs, and migrated 3 Power BI dashboards.
- Optimized a Dockerized Python AWS Batch pipeline (3 jobs: EC2, EBS, ELB) across 500+ accounts and regions; CUR-based filtering cut account-region-service combinations from 18,000 to 2,000 (89%) and weekly runtime from 8 hours to 2 hours (75%).
- Established governed sharing of billing data from 600+ AWS accounts with 10 consumer accounts using Lake Formation, S3 and AWS RAM (table grants, row-level filters, CloudFormation-automated IAM roles); eliminated manual approx. 200 GB transfers and approx. $1,000/month of query cost by moving querying to consumer accounts.
- Automated OCI Cost & Usage ingestion (Python, OCI SDK, Lambda, every 6 hours), replacing a daily manual download-and-mapping of business-unit compartment costs and eliminating human error; used DynamoDB file tracking, Pandas, Snappy Parquet and event-triggered Glue loads into Redshift.
- Delivered 8 Power BI dashboards for finance and business-unit leadership (AWS/OCI cost, budget forecasting, anomaly detection) on snowflake-schema data models with fact and dimension tables; incremental refresh in Dataflows Gen1 cut refresh time from 2 hours to 40 minutes (67%) on multi-million-record datasets.
- Implemented cost anomaly detection with CUR 2.0, Athena and Lambda (90-day Median + 2 MAD baseline, daily alerts); it flagged an approx. $1,000/day spend spike, leading to the responsible Lambda being stopped and avoiding an estimated $24,000 in charges.
- Implemented cross-region Redshift disaster recovery using scheduled snapshots and Lambda-orchestrated S3 Batch Replication; validated replication via object counts and job statuses with automated notifications.
- Developed a production Python MCP Server on AWS Lambda with Bedrock AgentCore Gateway and Okta authentication for AWS optimization discovery, S3 cost analysis and read-only Athena and Redshift queries through Claude Desktop.

Equiniti India, Chennai, India — Jul. 2022 – Dec. 2022
Process Executive

- Performed financial reconciliation, data validation and Excel reporting for operational accuracy.

Technical Skills

Languages, Libraries & Formats: Python, SQL, Pandas, NumPy, Parquet
Cloud (AWS & OCI): S3, Redshift, Glue, Glue Workflows, Athena, Lambda, Step Functions, Batch, DynamoDB, Lake Formation, CloudWatch, IAM, STS, Secrets Manager, OCI Object Storage
Data Engineering: ETL/ELT, Batch Pipelines, Data Warehousing, Data Integration, Incremental Processing, Data Validation, Data Quality, Data Governance, Disaster Recovery
Data Modeling & Databases: Dimensional Modeling, Snowflake Schema, Fact/Dimension Tables, PostgreSQL, MySQL
BI & Analytics: Power BI, Power Query, Dataflows Gen1, Incremental Refresh, Power BI Service, DAX (Basic)
DevOps & Tools: Infrastructure as Code (IaC), CloudFormation, YAML, Docker, Git, GitHub, Azure DevOps, ServiceNow
AI & Integration: Model Context Protocol (MCP), Amazon Bedrock AgentCore Gateway, Okta
Working Knowledge: Apache Spark, PySpark, Spark SQL, DataFrame API, Window Functions, Partitioning, Broadcast Joins, Microsoft Fabric, OneLake, Lakehouse, Dataflows Gen2, Fabric Data Pipelines

Certifications & Awards

AWS Certified Solutions Architect – Associate (SAA-C03) — Feb. 2024 – Feb. 2027
Microsoft Certified: Fabric Data Engineer Associate (DP-700) — Sep. 2026 – Sep. 2027
Microsoft Certified: Power BI Data Analyst Associate (PL-300) — Dec. 2024 – Dec. 2027
Grok Star, Datagrokr Analytics Pvt. Ltd. | *Outstanding work performance in Q1 2026* — Apr. 2026